

Lunar Lander


A photograph of two rockets launching from a launch pad. The rocket on the left is in the foreground, ascending vertically with a large plume of fire and smoke at its base. The rocket on the right is further away, also ascending with a smaller plume. The launch pad is visible in the foreground, and the sky is blue with scattered clouds.

Sarthak Khillon & Tim Stoddard




Introduction

- Lunar Lander v2 from the OpenAI Gym
- Land a simulated spaceship 🚀 on the moon!
- The terrain is randomly generated each round, but the designated landing spot is always flat
- Train lander to land itself, while maximizing score
- Many challenges!




Overview

- Each frame, lander can perform 4 actions
 - Do nothing
 - Fire left engine
 - Fire main (bottom) engine, -0.3 points per frame
 - Fire right engine
- Lander has infinite fuel
- Game is over when lander crashes (-100 points), or lands on the ground (+100 points)
 - Each leg on ground is +10 points
- Need at least 200 points to consider the challenge solved



Firing the main engine



Successful landing!



Value-based Learning (Q-learning)

- Learn a value function $V(s)$
 - Best sum of rewards up to that state.
 - “How good is this state?”
- Creates a policy π from V and evaluates best set of coefficients θ from multiple tests.
- Pros:
 - Independent of environment
- Cons:
 - Must discover policy by trial & error
 - Can have a large oscillation while training



```
valueIteration(S, A)
  V <- best value for each S

  for each state in S
    bestQ <- random value

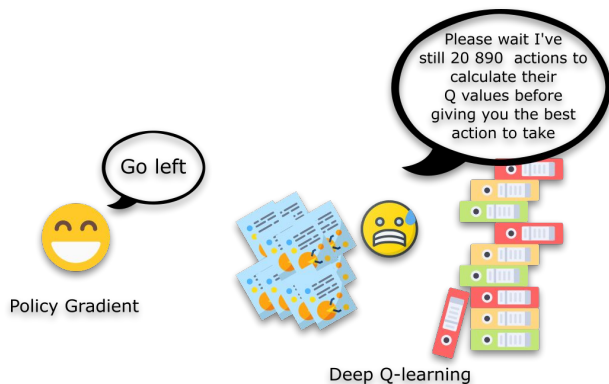
    for each action in A:
      currQ <- Q(state, action)
      update bestQ to currQ if better

  V[s] <- bestQ

  return V
```

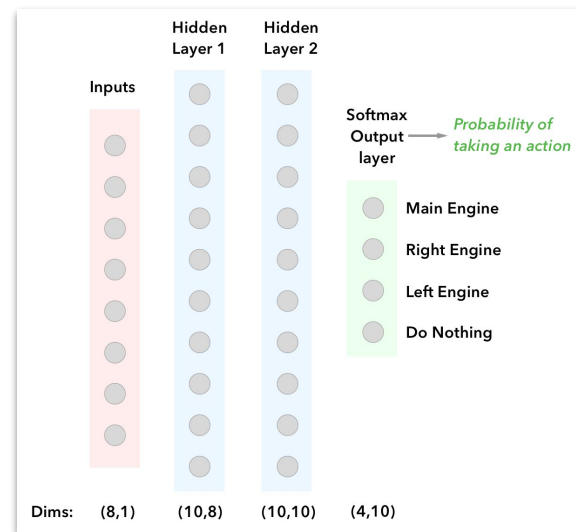
Policy-based Learning ★

- Goal is to find a policy $\pi(S_t \in S, A_t \in A)$ that maximizes reward
 - S is the state space, A is the action space
- Learn directly the policy function that maps a state to an action
 - Select actions without using a value function
- Converges faster than value-based (Q learning), but the policy evaluation step is more expensive
- How it applies to Lunar Lander
 - In each frame, the reward is a combination of how close the lander is to the landing pad and how close it is to zero speed
 - Closer to landing successfully = higher reward




Deep Deterministic Policy Gradient Network

- Approximates $\text{argmax}(a)$ on Q
- Advantages over a value-based approach
 - Better at convergence than value-based methods
 - Simply follow the gradient to find the best parameters
 - More effective in high dimensional action spaces, or when using continuous actions
- For loss, regular version uses softmax cross entropy, continuous version uses mean-squared error (MSE)
 - Adam Optimizer provides stochastic gradient descent



Xavier Initialization

- Initializing the network with the right weights is important for deep neural networks
 - If too small, input eventually drops to a really low value and can no longer be useful
 - If too large, it becomes so large that it becomes useless
- Need to make sure that the weights are in a reasonable range before we start training the network
- One good way is to assign the weights from a Gaussian distribution
 - This distribution would have zero mean and some finite variance



Algorithm Pseudocode

BEGIN

Continuous loop:

 Reset environment

 While not terminal condition:

 Choose and perform action.

 Environment returns next state + reward

 Update total reward for this episode

 Train: Fit θ values for policy

END

**Terminal state = crash, tilt, or landing.*

Libraries used

- numpy - fancy math things
- scipy - scientific computing
- tensorflow - dataflow programming
- gym - develop reinforcement learning algorithms




Demo!





Conclusion

- We created a Policy Gradient Network to solve the Lunar Lander v2 from the OpenAI Gym
- Policy gradients converge faster than value-based approaches, but are more computationally expensive
- Xavier initialization assigns network weights from Gaussian distribution
- Improvements - if we had more time
 - Hyperparameter tuning
 - Learning rates
 - Loss functions
 - Optimizer methods
- Where to go (future possibilities)
 - Based on problem, pick more general or more constrained network
 - Explore results from Actor-Critic and Q-learning approaches



References

1. [Deep Reinforcement Learning with Policy Gradients](#)
2. [Policy Gradients in a Nutshell](#)
3. [Tensorflow Deep Learning Projects](#)
4. [Deep Reinforcement Learning Demystified \(Episode 2\)](#)
5. [An introduction to Policy Gradients with Cartpole and Doom](#)
6. [Deep Q Network vs Policy Gradients](#)
7. [Reinforcement Learning Lecture, University of Washington](#)
8. [Understanding Xavier Initialization In Deep Neural Networks](#)

Images

- [Slide 2](#)
- [Slide 5](#)
- [Slide 6](#)
- [Slide 7](#)
- [Slide 9](#)
- [Slide 9](#)

